


Introduction to the Burrows-Wheeler Transform

Giovanni Manzini

1978 (?)



Space
Invaders
released

- David Wheeler conceives a data compression algorithm based on reversible transformation on the input text, but considers it too slow for practical use

1994



Mandela first
black South
Africa
president

- Mike Burrows improves the speed of the compressor. B&W co-author the technical report describing a “block sorting” lossless data compression algorithm.
- The algorithm splits the input in blocks and computes a reversible transformation that makes the text “more compressible”
- The transformation has been later called the Burrows-Wheeler transform.

1994

May 10, 1994

SRC Research
Report

124

**A Block-sorting Lossless
Data Compression Algorithm**

M. Burrows and D.J. Wheeler

digital

Systems Research Center
130 Lytton Avenue
Palo Alto, California 94301

The BWT

`swiss·miss·missing`

The BWT

swiss·miss·missing

Consider all rotations
of the input text

s wiss·miss·missin g
w iss·miss·missing s
i ss·miss·missings w
s s·miss·missingsw i
s ·miss·missingswi s
· miss·missingswis s
m iss·missingswiss ·
i ss·missingswiss· m
s s·missingswiss·m i
s ·missingswiss·mi s
· missingswiss·mis s
m issingswiss·miss ·
i ssingswiss·miss· m
s singswiss·miss·m i
s ingswiss·miss·mi s
i ngswiss·miss·mis s
n gswiss·miss·miss i
g swiss·miss·missi n

The BWT

swiss·miss·missing

Consider all rotations
of the input text

Sort them in
lexicographic order

· miss·missing**swiss** s
· missing**swiss**·mis s
g **swiss**·miss·missi n
i ng**swiss**·miss·mis s
i ss·miss·missing**s** w
i ss·missing**swiss**· m
i ssing**swiss**·miss· m
m iss·missing**swiss** ·
m issing**swiss**·miss ·
n g**swiss**·miss·miss i
s ·miss·missing**swi** s
s ·missing**swiss**·mi s
s ing**swiss**·miss·mi s
s s·miss·missing**sw** i
s s·missing**swiss**·m i
s sing**swiss**·miss·m i
s wiss·miss·missin g
w iss·miss·missing **s**

The BWT

swiss·miss·missing

Consider all rotations
of the input text

Sort them in
lexicographic order

Take the last character
of each rotation

ssnswmm··issssiigs

	L
· miss·missingswis	s
· missingswiss·mis	s
g swiss·miss·missi	n
i ngswiss·miss·mis	s
i ss·miss·missings	w
i ss·missingswiss·	m
i ssingswiss·miss·	m
m iss·missingswiss	·
m issingswiss·miss	·
n gswiss·miss·miss	i
s ·miss·missingswi	s
s ·missingswiss·mi	s
s ingswiss·miss·mi	s
s s·miss·missingsw	i
s s·missingswiss·m	i
s singswiss·miss·m	i
s wiss·miss·missin	g
w iss·miss·missing	s

The BWT

swiss·miss·missing

Consider all rotations
of the input text

Sort them in
lexicographic order

Take the last character
of each rotation

ssnswmm··issssiigs

F

L

·	miss·missingswis	s
·	missingswiss·mis	s
g	swiss·miss·missi	n
i	ngswiss·miss·mis	s
i	ss·miss·missings	w
i	ss·missingswiss·	m
i	ssingswiss·miss·	m
m	iss·missingswiss	·
m	issingswiss·miss	·
n	gswiss·miss·miss	i
s	·miss·missingswi	s
s	·missingswiss·mi	s
s	ingswiss·miss·mi	s
s	s·miss·missingsw	i
s	s·missingswiss·m	i
s	singswiss·miss·m	i
s	wiss·miss·missin	g
w	iss·miss·missing	s

Things to do next

1. Prove that given the transformed text we can retrieve T
2. Show that the transformed text is easy to compress

Fundamental observation

Every column of the matrix is a permutation of the input text (try to prove it!)




swiss · miss · miss**ing**

in F **s** is above **s** because
 ing ≤ wiss · · ·

in L **s** is in the row prefixed
 by **ing** hence is above **s**

F		L
·	miss · missing swiss	s
·	missing swiss · mis	s
g	swiss · miss · missi	n
i	ng swiss · miss · mis	s
i	ss · miss · missing s	w
i	ss · missing swiss ·	m
i	ssing swiss · miss ·	m
m	iss · missing swiss	·
m	issing swiss · miss	·
n	g swiss · miss · miss	i
s	· miss · missing swi	s
s	· missing swiss · mi	s
s	ing swiss · miss · mi	s
s	s · miss · missing sw	i
s	s · missing swiss · m	i
s	sing swiss · miss · m	i
s	wiss · miss · missin	g
w	iss · miss · missing	s

We can map each character in L to its image in F


F		L
·	miss·missing	swiss
·	missing	swiss·mis
g	swiss·miss·missi	n
i	ngswiss·miss·mis	s
i	ss·miss·missing	w
i	ss·missing	swiss·
i	ssing	swiss·miss·
m	iss·missing	swiss
m	issing	swiss·miss
n	gswiss·miss·miss	i
s	·miss·missing	swi
s	·missing	swiss·mi
s	ing	swiss·miss·mi
s	s·miss·missing	sw
s	s·missing	swiss·m
s	sing	swiss·miss·m
s	wiss·miss·missin	g
w	iss·miss·missing	s

We can map each character in L to its image in F

F		L
·	miss·missingswiss	s
·	missingswiss·mis	s
g	swiss·miss·missi	n
i	ngswiss·miss·mis	s
i	ss·miss·missings	w
i	ss·missingswiss·	m
i	ssingswiss·miss·	m
m	iss·missingswiss	·
m	issingswiss·miss	·
n	gswiss·miss·miss	i
s	·miss·missingswi	s
s	·missingswiss·mi	s
s	ingswiss·miss·mi	s
s	s·miss·missingsw	i
s	s·missingswiss·m	i
s	singswiss·miss·m	i
s	wiss·miss·missin	g
w	iss·miss·missing	s

We can map each
character in L to
its image in F

... even if we only
have L and F ...



Summing up

From L (the **BWT**) we can recover F

The relative order of the occurrences of any given character in F and L is the same

We can easily build a map telling us where each character in L is in F

We call this the **LF map**.


Using the LF map we
retrieve **T** right-to-left

T =

g

F
·
·
g
i
i
i
i
i
m
m
n
s
s
s
s
s
s
s
s
w


L
s
s
n
s
w
m
m
·
·
i
s
s
s
i
i
i
i
g
s



Using the LF map we
retrieve **T** right-to-left

T =


g



Using the LF map we
retrieve **T** right-to-left

T =


ng



Using the LF map we
retrieve T right-to-left

T =


ing



Using the LF map we
retrieve T right-to-left

T =


sing



Using the LF map we
retrieve T right-to-left

T =

ssing



What about compression?

Since BWT is a permutation of T, why should it be “easier” to compress?

Burrows and Wheeler observed that the BWT is usually “locally homogeneous” and suggested to compress it with Move-To-Front followed by an Order0 encoder (Huffman/Arithmetic Coding).

final char (<i>L</i>)	sorted rotations
a	n to decompress. It achieves compression
o	n to perform only comparisons to a depth
o	n transformation} This section describes
o	n transformation} We use the example and
o	n treats the right-hand side as the most
a	n tree for each 16 kbyte input block, enc
a	n tree in the output stream, then encodes
i	n turn, set $L[1]$ to be the
i	n turn, set $R[1]$ to the
o	n unusual data. Like the algorithm of Man
a	n use a single set of probabilities table
e	n using the positions of the suffixes in
i	n value at a given point in the vector R
e	n we present modifications that improve t
e	n when the block size is quite large. Ho
i	n which codes that have not been seen in
i	n with sch appear in the {\em same order
i	n with sch . In our exam
o	n with Huffman or arithmetic coding. Bri
o	n with figures given by Bell~\cite{bell}.


Figure 1: Example of sorted rotations. Twenty consecutive rotations from the sorted list of rotations of a version of this paper are shown, together with the final character of each rotation.

BWT vs $H_k(s)$ (1)

Let $s = \textit{ippississim}$,

$s^R = \textit{mississippi}$

$\text{BWT}(s^R) =$



i	m	i	s	s	i	s	s	i	p	p
i	p	p	i	m	i	s	s	i	s	s
i	s	s	i	p	p	i	m	i	s	s
i	s	s	i	s	s	i	p	p	i	m
m	i	s	s	i	s	s	i	p	p	i
p	i	m	i	s	s	i	s	i	p	p
p	p	i	m	i	s	s	i	s	i	s
s	i	p	p	i	m	i	s	s	i	s
s	i	s	s	i	p	p	i	m	i	s
s	s	i	p	p	i	m	i	s	s	i
s	s	i	s	s	i	p	p	i	m	i

BWT vs $H_k(s)$ (2)

Let $s = \textit{ippississim}$,

$s^R = \textit{mississippi}$

$\text{BWT}(s^R) =$

$$H_1(s) = (4/11) H_0(\textit{ps sm}) + (1/11) H_0(\textit{i}) \\ + (2/11) H_0(\textit{pi}) + (4/11) H_0(\textit{ss ii})$$

To compress up to $H_1(s)$ it suffices to compress each segment up to H_0

i	m	i	s	s	i	s	s	i	p	p
i	p	p	i	m	i	s	s	i	s	s
i	s	s	i	p	p	i	m	i	s	s
i	s	s	i	s	s	i	p	p	i	m
m	i	s	s	i	s	s	i	p	p	i
p	i	m	i	s	s	i	s	s	i	p
p	p	i	m	i	s	s	i	s	s	i
s	i	p	p	i	m	i	s	s	i	s
s	i	s	s	i	p	p	i	m	i	s
s	s	i	p	p	i	m	i	s	s	i
s	s	i	s	s	i	p	p	i	m	i

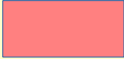
BWT vs $H_k(s)$ (3)

Let $s = \textit{ippississim}$,

$s^R = \textit{mississippi}$

$\text{BWT}(s^R) =$

i	mississipp	
ipp	mississ	
issipp	miss	
ississipp		
mississipp		
p	mississip	
pp	mississi	
sipp	missis	
sisipp	miss	
ssipp	missi	
ssissipp	imi	

To compress up to $H_2(s)$ it suffices to compress each segment  up to H_0

Summing up

To compress up to $H_k(s)$ it suffices to compress the corresponding partition of $BWT(s^R)$ up to H_0 (compare with PPM).

However:

- computing the partition is not easy
- which k should we choose?

It is possible to find an optimal partition in linear time, but there is a simpler alternative...

A closer look to MTF

ccbbaaabddddcccccc
11213112411411111

The integers are
in the range $[1, h]$
 h =alphabet size

Given s , $|\text{Gamma}(\text{MTF}(s))| \leq 2|s|H_0(s) + |s|$ only if in
the initial MTF list the symbols are in the same
order as in s , in our example: cbad.

If not, the penalty is at most $\log h$ bits per distinct
symbol, for any initial status of the MTF list
 $|\text{Gamma}(\text{MTF}(s))| \leq 2|s|H_0(s) + |s| + h \log h$ bits

Given three strings x, y, z

MTF($x \ y \ z$) differs from MTF(x) MTF(y) MTF(z) only for the status of the MTF list at the beginning of the encoding of y and z .

From the bound on the previous slide:

$$\begin{aligned} \text{Gamma}(\text{MTF}(x \ y \ z)) \leq & 2|x|H_0(x) + 2|y|H_0(y) + 2|z|H_0(z) \\ & + |x| + |y| + |z| + 3h \log h \end{aligned}$$

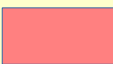
The same is true if we have more than three strings!

Since MTF has “little memory”, encoding the concatenation x_1, x_2, \dots, x_t produces an output bounded by:

$$\sum_i 2|x_i|H_0(x_i) + |x_i| + O(t \log h) \text{ bits}$$

This is precisely what we need for the BWT!

BWT(s^R) =

To compress up to $H_2(s)$ it suffices to compress each segment  up to H_0

i	m	i	s	s	i	s	s	i	p	p	
i	p	p	i	m	i	s	s	i	s	s	
i	s	s	i	p	p	i	m	i	s	s	
i	s	s	i	s	s	i	p	p	i	m	
m	i	s	s	i	s	s	i	p	p	i	
p	i	m	i	s	s	i	s	s	i	p	
p	p	i	m	i	s	s	i	s	s	i	
s	i	p	p	i	m	i	s	s	i	s	
s	i	s	s	i	p	p	i	m	i	s	
s	s	i	p	p	i	m	i	s	s	i	
s	s	i	s	s	i	p	p	i	m	i	


Main result

For any string s and order k :

$$|\text{Gamma}(\text{MTF}(\text{BWT}(s)))| \leq 2 |s| H_k(s) + |s| + O(h^{k+1} \log h) \text{ bits}$$

The bound can be improved replacing Gamma with another encoder, and applying Run Length Encoding (RLE) on the output of MTF

1997



First Harry
Potter novel
published

bzip2

- In 1997 Julian Seward released the bzip compression tool based on the BWT.
- The output of the BWT was compressed using Move-to-front, RLE, and Multiple Tables Huffman Coding
- It was highly optimized, and could be used as a drop-in replacement of gzip, both command line and library version

BWT as a compressor today

- Most of the mainstream compressors released in the last 20 years are based on LZ77 parsing, eg. LZMA, Snappy, Brotli, Zstd
- LZ77 has more “free parameters” and can offer a wide range of compression/speed trade-offs
- In BWT compression we cannot easily trade compression for speed

Compression results (1)

Size	Ratio %	C.MB/s	D.MB/s	Compressor (Binary 42% + Text 58%) Silesia.tar
48616057	22.9	1.07	77.11	LzTurbo 49
48758739	23.0	2.47	81.17	lzma 9
49517150	23.4	0.46	336.19	brotli 11d29
50861542	24.0	1.68	269.97	lzham 4
51720632	24.4	1.42	1239.95	LzTurbo 39
52715921	24.9	2.03	602.56	zstd 22
54596837	25.8	11.80	38.94	bzip2
58008992	27.4	7.96	853.20	zstd 15
59273940	28.0	59.48	1293.41	LzTurbo 32
59581397	28.1	33.48	416.81	brotli 5
60411647	28.5	45.64	798.97	zstd 9
60813803	28.7	1.60	2002.86	LzTurbo 29
64141404	30.3	162.02	1372.34	LzTurbo 31
64191258	30.3	65.28	416.81	brotli 4
64711652	30.5	0.22	325.27	zopfli
67624724	31.9	62.86	692.87	lzfse
67647204	31.9	9.99	316.72	zlib 9
68225985	32.2	24.46	313.67	zlib 6

bzip2 compresses well but it is relatively slow in compression and the slowest in decompression

Compression results (2)

Size	Ratio %	C.MB/s	D.MB/s	Compressor	Text log: NASA_access_log
11355945	5.5	0.86	320.68	LzTurbo 49	
11907661	5.8	0.99	2502.71	LzTurbo 39	
11960483	5.8	10.13	67.81	bzip2	
12236072	6.0	0.51	1022.47	brotli 11d29	
12617026	6.1	1.36	1348.32	zstd 22	
13598062	6.6	2.68	265.69	lzma 9	
13651218	6.7	1.33	880.25	lzham 4	
14661031	7.1	8.67	1819.99	zstd 15	
15041556	7.3	1.13	3732.63	LzTurbo 29	
16665926	8.1	78.89	1245.90	brotli 5	
17387746	8.5	117.98	1375.73	zstd 9	
18279979	8.9	187.64	2186.17	LzTurbo 32	
18654669	9.1	173.25	1227.89	brotli 4	
19085875	9.3	1.50	3527.36	lizard 49	
19545036	9.5	32.75	651.55	zlib 9	

for very compressible files bzip2 is more competitive in compression but still slow...